

Name:Kavya K
Usn:1NT18IS078
Subject:Big data lab

Use the Hadoop Mapreduce programming framework to compute with a program which will take the data from this .csv file and compute the following.

DATA SET:

	A	B	C	D	E	F
1	Kavya K	1NT18IS019	70	60	90	P
2	bhyravi	1NT18IS098	40	78	100	P
3	kshithija	1NT18IS084	50	78	99	P
4	pavithra	1NT18IS078	100	38	86	P
5	raheela	1NT18IS001	30	48	96	F
6	nischala	1NT18IS045	40	90	58	P
7	shamanth	1NT18IS002	40	78	48	P
8	vaishanavi	1NT18IS007	88	38	30	F
9	sannidhi	1NT18IS071	66	48	60	P
10	samruddhi	1NT18IS072	45	36	70	P
11	nikitha	1NT18IS073	67	65	20	F
12	cath	1NT18IS074	89	89	30	F
13	piyush	1NT18IS075	76	47	95	P
14	ahana	1NT18IS076	37	65	20	F
15	arushi	1NT18IS077	17	18	30	F
16	raghu	1NT18IS079	89	38	55	P
17	aravind	1NT18IS045	59	46	65	P
18	divya	1NT18IS046	76	55	70	P
19	bhavya	1NT18IS067	55	40	60	P
20	bhanu	1NT18IS043	55	63	43	P

- 1) Total number of students who have scored more than 60 in subject1

```

Student.java  Pass.java
1  package mypackone;
2
3  import java.io.IOException;
10
11  public class Student{
12
13  //MAPPER CODE
14
15  public static class Map extends MapReduceBase implements Mapper<LongWritable, Text, Text, IntWritable> {
16  private final static IntWritable one = new IntWritable(1);
17  private Text word = new Text();
18
19  public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter reporter)
20  throws IOException {
21  String myvalue = value.toString();
22  String[] transtokens = myvalue.split(",");
23
24  output.collect(new Text(transtokens[1]), new IntWritable(Integer.parseInt(transtokens[2]])); //{Mary, 1} //{had, 1} {little, 1} {lamb
25  }
26  }
27
28  //REDUCER CODE
29
30  public static class Reduce extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable> {
31  public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output,
32  Reporter reporter) throws IOException { //{little: {1,1}}
33  int limit = 60;
34  int val = Integer.MIN_VALUE;
35  while (values.hasNext()) {
36  if((val= values.next().get())>limit) {
37  output.collect(key, new IntWritable(val));
38  }
39  }
40  }
41  }
42
43  //DRIVER CODE
44  public static void main(String[] args) throws Exception {
45  JobConf conf = new JobConf(Student.class);
46  conf.setJobName("Transaction Counting");
47  conf.setOutputKeyClass(Text.class);
48  conf.setOutputValueClass(IntWritable.class);
49  conf.setMapperClass(Map.class);
50  conf.setCombinerClass(Reduce.class);
51  conf.setReducerClass(Reduce.class);
52  conf.setInputFormat(TextInputFormat.class);
53  conf.setOutputFormat(TextOutputFormat.class); // hadoop jar jarname classpath inputfolder outputfolder
54  FileInputFormat.setInputPaths(conf, new Path(args[0]));

```

```

37         output.collect(key, new IntWritable(val));
38     }
39 }
40 }
41 }
42
43 //DRIVER CODE
44 public static void main(String[] args) throws Exception {
45     JobConf conf = new JobConf(Student.class);
46     conf.setJobName("Transaction Counting");
47     conf.setOutputKeyClass(Text.class);
48     conf.setOutputValueClass(IntWritable.class);
49     conf.setMapperClass(Map.class);
50     conf.setCombinerClass(Reduce.class);
51     conf.setReducerClass(Reduce.class);
52     conf.setInputFormat(TextInputFormat.class);
53     conf.setOutputFormat(TextOutputFormat.class); // hadoop jar jarname classpath inputfolder outputfolder
54     FileInputFormat.setInputPaths(conf, new Path(args[0]));
55     FileOutputFormat.setOutputPath(conf, new Path(args[1]));
56     JobClient.runJob(conf);
57 }
58 }
59 }
60

```

```

[000] password for kavya:
hadoop@kavya-VirtualBox: /home/kavya/Desktop$ cd /usr/local/hadoop/sbin
hadoop@kavya-VirtualBox: /usr/local/hadoop/sbin$ ./stop-all.sh
WARNING: Stopping all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: Use CTRL-C to abort.
Stopping namenodes on [localhost]
Stopping datanodes
Stopping secondary namenodes [kavya-VirtualBox]
Stopping nodemanagers
Stopping resourcemanager
hadoop@kavya-VirtualBox: /usr/local/hadoop/sbin$ ./start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [kavya-VirtualBox]
Starting resourcemanager
Starting nodemanagers
hadoop@kavya-VirtualBox: /usr/local/hadoop/sbin$ source /etc/profile.d/hadoop_java.sh
hadoop@kavya-VirtualBox: /usr/local/hadoop/sbin$ hdfs dfs -put /home/kavya/Downloads/MARK.csv /du
hadoop@kavya-VirtualBox: /usr/local/hadoop/sbin$ hadoop jar /home/kavya/Desktop/marks.jar /du /aru
2021-07-15 19:30:35,695 INFO impl.MetricsConfig: loaded properties from hadoop-metrics2.properties
2021-07-15 19:30:36,033 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2021-07-15 19:30:36,033 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2021-07-15 19:30:36,222 WARN impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2021-07-15 19:30:37,000 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface
execute your application with ToolRunner to remedy this.
2021-07-15 19:30:37,528 INFO mapred.FileInputFormat: Total input files to process : 1
2021-07-15 19:30:37,796 INFO mapreduce.JobSubmitter: number of splits:1
2021-07-15 19:30:38,807 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1486643880_0001
2021-07-15 19:30:38,814 INFO mapreduce.JobSubmitter: Executing with tokens: []
2021-07-15 19:30:39,497 INFO mapreduce.Job: The url to track the job: http://localhost:8080/

```

```
hadoop@kavya-VirtualBox: /usr/local/hadoop/sbin

Spilled Records=18
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=140
Total committed heap usage (bytes)=274866176

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
Bytes Read=617
File Output Format Counters
Bytes Written=127

hadoop@kavya-VirtualBox:/usr/local/hadoop/sbin$ hdfs dfs -ls /aru
Found 2 items
-rw-r--r-- 1 hadoop supergroup          0 2021-07-15 19:30 /aru/_SUCCESS
-rw-r--r-- 1 hadoop supergroup       127 2021-07-15 19:30 /aru/part-00000

hadoop@kavya-VirtualBox:/usr/local/hadoop/sbin$ hdfs dfs -cat /aru/part*/
1NT18IS007      88
1NT18IS019      70
1NT18IS046      76
1NT18IS071      66
1NT18IS073      67
1NT18IS074      89
1NT18IS075      76
1NT18IS078     100
1NT18IS079      89
hadoop@kavya-VirtualBox:/usr/local/hadoop/sbin$
```

2) Total number of students who have passed in all subjects.

```
Student.java Pass.java
1 package mypackone;
2
3
4
5 import java.io.IOException;
6 import java.util.*;
7 import org.apache.hadoop.fs.Path;
8 import org.apache.hadoop.conf.*;
9 import org.apache.hadoop.io.*;
10 import org.apache.hadoop.mapred.*;
11 import org.apache.hadoop.util.*;
12
13 public class Pass{
14
15 //MAPPER CODE
16
17 public static class Map extends MapReduceBase implements Mapper<LongWritable, Text, Text, IntWritable> {
18 private final static IntWritable one = new IntWritable(1);
19 private Text word = new Text();
20
21 public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter reporter)
22 throws IOException {
23 String myvalue = value.toString();
24 String[] transtokens = myvalue.split(",");
25 if(transtokens [5].equals("P")) {
```

```
26     output.collect(new Text(transtokens[1]), one); //{Mary, 1} //{had, 1} {little, 1} {lamb,1} {little, 1}
27 }
28 }
29
30 }
31
32 //REDUCER CODE
33 public static class Reduce extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable> {
34 public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output,
35 Reporter reporter) throws IOException { //{little: {1,1}}
36 int count = 0;
37
38 while (values.hasNext()) {
39
40 count += values.next().get();
41
42 }
43 output.collect(key, new IntWritable(count));
44 }
45 }
46
47 //DRIVER CODE
48 public static void main(String[] args) throws Exception {
```

```
Student.java Pass.java
40 count += values.next().get();
41
42 }
43 output.collect(key, new IntWritable(count));
44 }
45 }
46
47 //DRIVER CODE
48 public static void main(String[] args) throws Exception {
49 JobConf conf = new JobConf(Pass.class);
50 conf.setJobName("Transaction Counting");
51 conf.setOutputKeyClass(Text.class);
52 conf.setOutputValueClass(IntWritable.class);
53 conf.setMapperClass(Map.class);
54 conf.setCombinerClass(Reduce.class);
55 conf.setReducerClass(Reduce.class);
56 conf.setInputFormat(TextInputFormat.class);
57 conf.setOutputFormat(TextOutputFormat.class); // hadoop jar jarname classpath inputfolder outputfolder
58 FileInputFormat.setInputPaths(conf, new Path(args[0]));
59 FileOutputFormat.setOutputPath(conf, new Path(args[1]));
60 JobClient.runJob(conf);
61 }
62 }
63
```

```

hadoop@kavya-VirtualBox: /usr/local/hadoop/sbin$ hadoop jar /home/kavya/Desktop/any.jar /cs /number
2021-07-15 20:27:38,635 INFO impl.MetricsConfig: loaded properties from hadoop-metrics2.properties
2021-07-15 20:27:39,325 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2021-07-15 20:27:39,325 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2021-07-15 20:27:39,435 WARN impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2021-07-15 20:27:40,303 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and
execute your application with ToolRunner to remedy this.
2021-07-15 20:27:40,988 INFO mapred.FileInputFormat: Total input files to process : 1
2021-07-15 20:27:41,758 INFO mapreduce.JobSubmitter: number of splits:1
2021-07-15 20:27:42,868 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local27726382_0001
2021-07-15 20:27:42,871 INFO mapreduce.JobSubmitter: Executing with tokens: []
2021-07-15 20:27:43,622 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2021-07-15 20:27:43,630 INFO mapreduce.Job: Running job: job_local27726382_0001
2021-07-15 20:27:43,654 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2021-07-15 20:27:43,706 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapred.FileOutputCommitter
2021-07-15 20:27:43,757 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2021-07-15 20:27:43,764 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ign
ore cleanup failures: false
2021-07-15 20:27:44,106 INFO mapred.LocalJobRunner: Waiting for map tasks
2021-07-15 20:27:44,130 INFO mapred.LocalJobRunner: Starting task: attempt_local27726382_0001_m_000000_0
2021-07-15 20:27:44,302 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2021-07-15 20:27:44,302 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ign
ore cleanup failures: false
2021-07-15 20:27:44,608 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2021-07-15 20:27:44,710 INFO mapred.MapTask: Processing split: hdfs://localhost:9000/cs/0+617
2021-07-15 20:27:44,730 INFO mapreduce.Job: Job job_local27726382_0001 running in uber mode : false
2021-07-15 20:27:44,732 INFO mapreduce.Job: map 0% reduce 0%
2021-07-15 20:27:44,886 INFO mapred.MapTask: numReduceTasks: 1
2021-07-15 20:27:45,710 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
2021-07-15 20:27:45,710 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100

```

Activities

Terminal

Jul 15 20:29

hadoop@kavya-VirtualBox: /usr/local/hadoop/sbin

Spilled Records=26

Shuffled Maps =1

Failed Shuffles=0

Merged Map outputs=1

GC time elapsed (ms)=185

Total committed heap usage (bytes)=274866176

Shuffle Errors

BAD_ID=0

CONNECTION=0

IO_ERROR=0

WRONG_LENGTH=0

WRONG_MAP=0

WRONG_REDUCE=0

File Input Format Counters

Bytes Read=617

File Output Format Counters

Bytes Written=169

hadoop@kavya-VirtualBox: /usr/local/hadoop/sbin\$ hdfs dfs -cat /number/part*

1NT18IS002

1

1NT18IS019

1

1NT18IS043

1

1NT18IS045

2

1NT18IS046

1

1NT18IS067

1

1NT18IS071

1

1NT18IS072

1

1NT18IS075

1

1NT18IS078

1

1NT18IS079

1

1NT18IS084

1

1NT18IS098

1

hadoop@kavya-VirtualBox: /usr/local/hadoop/sbin\$